

Governance in Hospital Management: The NHS Patient Flow and Bed Management Case

Pradyumna Reddy Dereddi, Nishanth Reddy Dereddi, Sireesha Rodda*

School of Computer Science and Engineering, GITAM deemed to be University,
Visakhapatnam, India

*Corresponding author: srodde@gitam.edu

Abstract

The usage of Artificial Intelligence (AI) has offered premium quality of service at several organizations. This paper explores the case of hospital management of use of artificial intelligence in patient flow and bed management. Observations are drawn from documented case study of National Health Service in UK, Birmingham and Leeds Teaching Hospitals on how factors such as algorithmic transparency, regulatory frameworks, accountability and equity considerations shape the domain of operational hospital management based on AI. The case study shows how artificial intelligence can be smartly used to improve on reducing the duration of stay also improving bed utilization. An adaptive governance framework of public health systems, which integrates artificial intelligence into hospital operations is analyzed. The framework lays ground to reduce bed occupancy delays and improve the efficiency in operations, resulting in equitable and high-quality patient care.

Keywords: Artificial Intelligence, AI Governance, hospital management, patient management

1. Introduction

The foray of Artificial Intelligence into medicine has brought fundamental transformation in more ways than one. From diagnostic imaging to surgical scheduling, AI systems have now become indispensable across the clinical and operational layers of hospital management. Among the operational applications, managing patient flow and bed allocation account for significant challenges to be addressed by AI.

The National Health Service (NHS) in the United Kingdom represents one of the most ambitious organizations which has well documented the attempts to integrate AI into hospital management at large scale. Ageing population, rising demand and limited

resources pose serious challenges to NHS. It invests substantially into AI tools designed to predict patient admissions, to anticipate discharge readiness, and to optimise the bed allocation. All these efforts are in alignment with the long-term goals of NHS towards digital transformation. (NHS England, 2019).

The deployment of AI in hospital management is not straightforward technical action but wrought in challenges like clinical accountability, institutional trust, constraints on data use, and ethical safeguarding. This paper examines the NHS's experience to draw lessons for AI governance in public health systems.

2. The NHS Case: AI for Optimising Patient Flow and Bed Management

University Hospitals Birmingham (UHB) and Leeds Teaching Hospitals have been among the most prominent NHS trusts to deploy AI-enabled patient flow systems. An AI platform built on Microsoft Azure and Epic Electronic Health records (HER) generates real-time likelihood scores for every inpatient, drawing on clinical indicators, historical length-of-stay patterns, and care pathway data. Bed managers and ward teams use the provided dashboards to plan patient movements 24 to 48 hours in advance, shifting from reactive to predictive operational management (University Hospitals Birmingham NHS Foundation Trust, 2021).

Reported outcomes have been promising. UHB documented a reduction in average length of stay of approximately 0.5 days per patient — a significant aggregate efficiency gain across more than 100,000 annual admissions. Leeds reported improvements in emergency department throughput as upstream bed availability improved, reducing the logjam that historically delayed patient admissions from emergency to inpatient settings. Nationally, NHS England's Patient Flow Intelligence Tool (PFIT) has extended similar predictive analytics to system-level commissioners, enabling demand modelling across regional health economies (NHS England, 2021).

The aforementioned gains are heavily dependent on data quality, staff adoption rates, and organisational culture (Nuffield Trust, 2022). This evidence prompts the need for AI governance adding value to hospital management.

3. Governance Analysis

3.1 Regulatory Framework

The regulatory governance of AI in NHS hospital management spans multiple bodies. The Medicines and Healthcare products Regulatory Agency (MHRA) oversees AI tools classified as medical devices, while NHS Digital enforces data security standards, and the Care Quality Commission (CQC) assesses the safety and quality of care in NHS settings. For AI tools that support operational decisions without directly informing clinical treatment — as is largely the case with bed management dashboards — regulatory requirements are less stringent, falling primarily under the Data Protection Act 2018 and NHS Digital's Data Security and Protection Toolkit (MHRA, 2021).

This regulatory landscape, while comprehensive in principle, is fragmented in practice. NHS trusts deploying AI tools must navigate overlapping mandates from several regulatory bodies, creating compliance uncertainty. The NHS AI and Digital Regulations Service, a joint initiative between MHRA, CQC, and NHS Digital, has begun to consolidate guidance, but trusts report that clearer, unified regulatory pathways are needed (Nuffield Trust, 2022).

3.2 Transparency and Explainability

The machine learning models underpinning NHS bed management tools — including gradient boosting algorithms and neural networks — are powerful but often opaque. A bed manager receiving a discharge likelihood score has no inherent visibility into which data features drove the prediction. This opacity creates clinical governance risks: professionals may act on algorithmic recommendations without being able to critically evaluate their basis.

NHS England has encouraged trusts to demand explainability features from AI vendors, such as SHAP (SHapley Additive exPlanations) values that surface the factors most influencing a given prediction (NHS England, 2021). UHB has incorporated feature attribution displays into its dashboards. However, implementation of such features remains inconsistent across NHS trusts, and their utility depends on staff AI literacy — a competency that is unevenly distributed across the NHS workforce.

3.3 Equity and Bias

AI systems trained on NHS data inherit the biases present in historical clinical records. Because Black, Asian, and Minority Ethnic (BAME) patients have historically experienced differential care pathways — including longer waits, delayed diagnoses, and

disparate discharge planning — AI models trained on this data risk encoding and amplifying these inequities. The NHS Race and Health Observatory has identified algorithmic bias as a priority concern in NHS AI deployments (NHS Race and Health Observatory, 2021).

Critically, there is no current requirement for mandatory equity audits of operational AI tools in the NHS — only of AI medical devices. This gap means that bed management algorithms may operate without ever being scrutinised for differential performance across patient demographic groups. Addressing this requires embedding equity metrics in AI procurement standards and establishing regular disaggregated performance reporting.

3.4 Accountability and Staff Engagement

Accountability in AI-assisted hospital management remains a contested question. NHS trust governance places ultimate liability with the employing organisation and the responsible clinician, but the presence of an algorithm in the decision chain creates ambiguity about where professional judgment ends and algorithmic influence begins. NHS England's AI assurance framework requires that human oversight is maintained for all consequential decisions, and that AI tools are used as decision support rather than autonomous decision-makers (NHS England, 2021).

Staff engagement is equally critical. Clinicians and bed managers who distrust or misunderstand the AI system are unlikely to use it effectively, nullifying its operational benefits. The UHB deployment succeeded in part because it involved frontline staff in co-design from the outset, providing training and feedback mechanisms that built confidence in the system's outputs. Governance frameworks must therefore treat staff engagement not as a communication afterthought but as a structural component of AI deployment.

4. Proposed Governance Framework

Based on the analysis, this paper proposes four core principles for governing AI in public hospital management. Proportionate regulation is much needed. A tiered approach, which demarcates clinical AI and administrative AI, can enable the resources to be allocated resources more effectively. Mandatory equity auditing, which should require all AI tools in the NHS settings, to be subject to equity audits examining differential performance across patient demographic groups. Algorithmic transparency must be implemented as a procurement condition. NHS trusts should require vendors to provide explainability features,

audit access, and algorithmic documentation as conditions of contract. Proprietary black-box systems, that do not encourage trust among clinicians, or regulators should not be eligible for NHS deployment. Fourth, living governance which requires AI systems must be subject to continuous monitoring against performance and equity benchmarks, with governance structures empowered to mandate modifications or suspend deployment when harms are identified. This adaptive principle recognises that AI systems evolve as underlying data distributions shift, and that governance must evolve with them. The detailed framework is presented in Figure 1.




Figure 1: AI Governance Framework-System Architecture

5. AI Techniques in Patient Flow and Bed Management

This section outlines the AI techniques that could be deployed for various modules of NHS patient flow and bed management system.

5.1 Predictive Modelling

Ensemble boosting techniques such as XGBoost, LightGBM maybe utilised to generate discharge likelihood scores from structured electronic health record (EHR) data, including clinical observations, diagnostic results, medication records, and historical length-of-stay patterns. These ensemble methods rely on the combined decisions of multiple base models to minimise prediction error, achieving high accuracy on tabular clinical data. The base models are usually simpler classifications models which predict situations like: “will this patient be discharged within 24 hours?”

5.2 Deep Learning and Natural Language Processing

More recent deployments have incorporated deep learning techniques to process the sequential and unstructured dimensions of clinical data. Recurrent Neural Networks (RNNs) and Long Short-Term Memory (LSTM) architectures are particularly suited to time-series clinical data, which include analysing patient’s vital signs, medication changes, or physiological observations over time to detect patterns indicative of imminent discharge readiness or deterioration. Transformer-based architectures, originally developed for natural language tasks, are increasingly applied to process free-text clinical notes within NHS EHR systems. Natural Language Processing (NLP) techniques — including Named Entity Recognition (NER) and text classification — extract clinically relevant information from nursing notes and medical records, enabling the AI system to identify discharge barriers such as “awaiting social care assessment” or “pending occupational therapy review” that would not appear in structured data fields.

5.3 Explainability Techniques

Establishing clinical trust when deploying AI models is crucial. Attaching explainability to AI models is imperative for algorithmic transparency discussed in Section 3.2, explainability techniques form a critical component of the NHS AI stack. SHAP (SHapley Additive exPlanations) assigns a contribution value to each feature for a specific prediction. In practice, a bed manager reviewing a low discharge likelihood score can see that the prediction is driven primarily by an outstanding physiotherapy assessment and an abnormal blood result, rather than demographic proxies. LIME (Local Interpretable Model-agnostic Explanations) offers a complementary approach, fitting a locally interpretable linear model around individual predictions to approximate the behaviour of more complex architectures. Both techniques are model-agnostic, and can be applied to any predictive model regardless of its internal architecture.

5.4 Demand Forecasting and Optimisation

Time series models such as ARIMA can be used to predict admission volumes and enable proactive capacity planning. These models underpin the NHS Patient Flow Intelligence Tool (PFIT), which operates at regional and national scale. Once demand is forecast, optimisation techniques translate predictions into actionable allocation decisions. Reinforcement Learning based agents can learn dynamic bed allocation policies through simulated interactions with the hospital environment, adapting in real time to changing ward conditions.

5.5 Data Infrastructure Techniques

Federated Learning enables AI models to be trained across multiple NHS trusts by sharing model gradients rather than raw patient data, preserving privacy while allowing the model to learn from a broader and more representative population. This is particularly valuable for addressing the bias risks identified in Section 3.3. Transfer Learning complements this approach by allowing pre-trained models, developed on large, multi-trust datasets, to be fine-tuned on a specific trust's local data, reducing the volume of labelled training data required for effective deployment at smaller or less digitally mature trusts.

6. Conclusions

The NHS's deployment of AI for patient flow and bed management demonstrates that AI can meaningfully improve hospital operational efficiency, but only when accompanied by governance frameworks that address transparency, equity, accountability, and staff engagement.

Effective AI governance requires that administrators, clinicians, policymakers, and technology developers work in unison to ensure that AI serves the foundational values of public healthcare: safety, equity, and the primacy of patient welfare. The NHS's experience offers both encouragement and caution for health systems worldwide embarking on this transformation.

References

- Medicines and Healthcare products Regulatory Agency [MHRA]. (2021). Software and AI as a medical device change programme. UK Government. <https://www.gov.uk/government/publications/software-and-ai-as-a-medical-device-change-programme>

- NHS England. (2019). The NHS long term plan. <https://www.longtermplan.nhs.uk>
- NHS England. (2021). A buyer's guide to AI in health and care. NHS Transformation Directorate. <https://transform.england.nhs.uk/ai-lab/>
- NHS Race and Health Observatory. (2021). Ethnic health inequalities and the NHS. <https://www.nhsrho.org>
- Nuffield Trust. (2022). Artificial intelligence in the NHS: An evaluation of early deployments. <https://www.nuffieldtrust.org.uk>
- Obermeyer, Z., Powers, B., Vogeli, C., & Mullainathan, S. (2019). Dissecting racial bias in an algorithm used to manage the health of populations. *Science*, 366(6464), 447–453. <https://doi.org/10.1126/science.aax2342>
- Topol, E. J. (2019). High-performance medicine: The convergence of human and artificial intelligence. *Nature Medicine*, 25(1), 44–56. <https://doi.org/10.1038/s41591-018-0300-7>
- University Hospitals Birmingham NHS Foundation Trust. (2021). AI-powered patient flow: Evaluation report 2020–2021 [Internal Report]. UHB Digital & Informatics.
- World Health Organization [WHO]. (2021). Ethics and governance of artificial intelligence for health. WHO Press. <https://www.who.int/publications/i/item/9789240029200>